

Project Goal

Develop a campus-relevant, open-source AI-augmented SOC prototype using **synthetic/anonymized CSUSB security datasets**.

Focus: log summarization, alert triage, and report generation using LLMs + agent frameworks.

1. Project Teams and Roles

- **Team 1 — SOC Infrastructure:** Deploy Wazuh + Elastic/OpenSearch, TheHive + Cortex, Zeek/Suricata, and Docker Compose
- **Team 2 — Log Summarization & LLM Integration:** LibreLog, LogBatcher, LILAC, LLaMA/Mistral models
- **Team 3 — Alert Triage & Orchestration:** LangChain agent, Chroma vector DB, FastAPI microservices
- **Team 4 — Report Generation:** AGIR, GenDFIR, AttackGen
- **Team 5 — Testing & Evaluation:** Metrics, benchmarks, dashboards, documentation

Each team has primary responsibilities but should collaborate closely for integration.

2. Required Tools and Resources

SOC Infrastructure:

- Wazuh (IDS/SIEM)
- Elastic/OpenSearch (log storage & querying)
- TheHive + Cortex (incident response & ticketing)
- Zeek / Suricata (network traffic analysis)
- Docker Compose (deployment & reproducibility)

AI/LLM Tools:

- LibreLog, LogBatcher, LILAC (log parsing)
- LLaMA 3 8B / Mistral-7B (local LLM models)
- LangChain or LangGraph (agent orchestration)
- ChromaDB (vector database)
- FastAPI (microservices)

Datasets:

- Public/academic: CICIDS 2017, UNSW-NB15, NSL-KDD, Kyoto 2006+, LogHub-2.0
 - Campus-specific: CSUSB synthetic or anonymized logs (after approval)
-

3. Step-by-Step 3-Month Roadmap

Month 1 — SOC Setup

Objectives: Deploy basic SOC infrastructure and ingest sample datasets.

Tasks:

- 1. Install & configure:**
 - Wazuh manager & agents
 - Elastic/OpenSearch with dashboards
 - TheHive + Cortex for incident/ticket management
 - Zeek / Suricata for network traffic monitoring
- 2. Dockerize** each component for reproducibility
- 3. Replay sample datasets:**
 - CICIDS 2017 network traffic
 - LogHub system logs
- 4. Verify ingestion & dashboards:**
 - Ensure logs appear in Elastic
 - Test alert generation via Zeek/Suricata

5. Documentation:

- Record setup steps, config files, Docker Compose structure
-

Month 2 — AI Augmentation

Objectives: Add AI/LLM modules for log summarization, alert triage, and report generation.

Tasks:

1. Log Summarization:

- Integrate LibreLog / LogBatcher / LILAC
- Process logs from Month 1
- Output: concise summaries for each log batch

2. Alert Triage:

- Implement LangChain agent
- Connect to Chroma vector DB
- Automate classification & prioritization of alerts
- Output: reduced false positives, alerts with context

3. Report Generation:

- Implement AGIR / GenDFIR / AttackGen
- Convert alerts and summaries into human-readable reports
- Output: ready-to-read incident reports

4. Integration:

- Connect log summarization → alert triage → report generation pipeline

5. Testing:

- Run small-scale test with CICIDS + synthetic CSUSB logs

6. Documentation:

- Include API endpoints, configuration, and usage instructions

Month 3 — Testing, Benchmarking, and Demo

Objectives: Evaluate system performance, optimize, and prepare deliverables.

Tasks:

1. Evaluation Metrics:

- Log summarization: compare summaries with human-annotated logs → accuracy or BERTScore
- Alert triage: false positives, precision/recall, F1-score
- Report generation: time saved per incident, completeness of details

2. Benchmarking:

- Run both baseline SOC (without AI) and AI-augmented SOC
- Compare metrics for measurable improvement

3. CSUSB Scenario Simulation:

- Use synthetic/anonymized CSUSB traffic for realistic evaluation

4. Demo Preparation:

- Dockerized AI-SOC ready for presentation
- Slide deck summarizing methodology, results, and learning outcomes

5. Final Documentation:

- Detailed report of setup, AI integration, testing, and results
- GitHub repository with Dockerfiles, notebooks, datasets (where allowed)

4. Deliverables

- Fully **Dockerized AI-augmented SOC**
- Benchmarked **performance metrics**

- **Incident reports** generated by the system

- **Final presentation slides**

GitHub repository with scripts, configs, and documentation